

## 3.2 Linear Regression – Stochastic Gradient Descent

Lecture based on “Dive into Deep Learning” <http://D2L.AI> (Zhang et al., 2020)

Prof. Dr. Christoph Lippert


Digital Health & Machine Learning

The **loss function** quantifies the distance between the **real** and **predicted** value of the target.

$$l^{(i)}(\mathbf{w}, b) = \frac{1}{2} \left( \mathbf{w}^\top \mathbf{x}^{(i)} + b - y^{(i)} \right)^2.$$

$$L(\mathbf{w}, b) = \frac{1}{n} \sum_{i=1}^n l^{(i)}(\mathbf{w}, b)$$

$$(\mathbf{w}^*, b^*) = \underset{\mathbf{w}, b}{\operatorname{argmin}} L(\mathbf{w}, b)$$



- The **OLS** can be determined analytically.
- An alternative algorithm is **gradient descent**.

## Optimization

### Gradient descent


$$(\mathbf{w}^*, b^*) = \operatorname{argmin}_{\mathbf{w}, b} L(\mathbf{w}, b)$$

$-\nabla L(\mathbf{w}, b)$  is the direction of **steepest descent**.

**Gradient descent** update rule:

$$(\mathbf{w}, b) \leftarrow (\mathbf{w}, b) - \eta \nabla L(\mathbf{w}, b),$$

for a small **stepsize**  $\eta > 0$  (e.g.,  $10^{-3}$ ).



### Note:

- $O(n)$  cost for gradient computation.
- Can be very slow for some problems (oscillation).
- Given sufficient time, gradient descent will find the minimum
- The loss of linear regression loss is (strictly) **convex**.

## Stochastic Gradient Descent

- Initialize the model parameters (randomly)
- Iterate:
  - i Uniformly sample a mini-batch  $\mathcal{B}$  of training examples.
  - ii Compute the gradient of the average loss on the mini batch.
  - iii update.

$$(\mathbf{w}, b) \leftarrow (\mathbf{w}, b) - \frac{\eta}{|\mathcal{B}|} \sum_{i \in \mathcal{B}} \partial_{(\mathbf{w}, b)} l^{(i)}(\mathbf{w}, b)$$

### Note:


- The batch size and learning rate  $\eta > 0$  are **hyper-parameters** that have to be **tuned** by the user.
- Avoids  $O(n)$  cost for gradient computation.
- Can still oscillate.
- Never reaches the exact optimum.

$$(\mathbf{w}, b) \leftarrow (\mathbf{w}, b) - \frac{\eta}{|\mathcal{B}|} \sum_{i \in \mathcal{B}} \partial_{(\mathbf{w}, b)} l^{(i)}(\mathbf{w}, b)$$

For quadratic losses and linear functions we can compute the updates explicitly:

$$\begin{aligned} \mathbf{w} &\leftarrow \mathbf{w} - \frac{\eta}{|\mathcal{B}|} \sum_{i \in \mathcal{B}} \partial_{\mathbf{w}} l^{(i)}(\mathbf{w}, b) &= \mathbf{w} - \frac{\eta}{|\mathcal{B}|} \sum_{i \in \mathcal{B}} \mathbf{x}^{(i)} \left( \mathbf{w}^\top \mathbf{x}^{(i)} + b - y^{(i)} \right), \\ b &\leftarrow b - \frac{\eta}{|\mathcal{B}|} \sum_{i \in \mathcal{B}} \partial_b l^{(i)}(\mathbf{w}, b) &= b - \frac{\eta}{|\mathcal{B}|} \sum_{i \in \mathcal{B}} \left( \mathbf{w}^\top \mathbf{x}^{(i)} + b - y^{(i)} \right). \end{aligned}$$

## A Single-Layer Neural Network.




Linear models have a single neuron.

- $d$  inputs  $x_1, x_2, \dots, x_d$
- 1 output

All inputs are connected to all outputs.

- **fully-connected layer** or **dense layer**.



## Summary

- Linear regression with a squared loss has an analytic solution.
- Gradient descent is an iterative algorithm to minimize a differentiable loss function
  - Each step of gradient descent requires a pass over the whole data set.
  - Gradient descent can be very slow (e.g. oscillation)
- Stochastic gradient descent speeds up gradient calculation, by approximating the gradient on a small random subsample of the training data.
  - Requires hyperparameter tuning to perform well
  - Goto optimization algorithm in deep learning
- Linear models are single layer neural networks.