

## 4.2. Generalization and Regularization

Lecture based on “Dive into Deep Learning” <http://D2L.AI> (Zhang et al., 2020)  
and C.M. Bishop, Pattern Recognition and Machine Learning

Prof. Dr. Christoph Lippert

Digital Health & Machine Learning

## Training and generalization error

- *Training error*: The error calculated on the training data set
- *Generalization error*: The expectation of our model's error on additional data points drawn from the same underlying data distribution.
- Problem: *we can never calculate the generalization error exactly*
- → apply model to an independent test set, withheld from training

### Example (Coin Toss)

- Dataset  $\{0, 1, 1, 1, 0, 1\}$  → Predict the *majority class* (here: 1) with error of only  $\frac{1}{3}$ .
- With more samples it would go to  $\frac{1}{2}$ .

**Generalization is the fundamental problem in machine learning.**

## Drawing training and validation samples

In **supervised learning**, we usually assume that both the training data and the test data are drawn *independently* from *identical* distributions (i.i.d.).

- Sampling process has no memory
- 2nd and 3rd samples are no more correlated than 2nd and  $n$ th sample

### Example

Covid-19 mortality risk predictor on data collected from patients from Charité Berlin, and apply it on patients from Mt. Sinai.

### Example

Face recognition trained only on *students* and then applied in an *elderly home*

**Goal: Find a function that fits training set well, but generalizes well on unseen data**


## Underfitting or overfitting

**Aim:** Model, where training error and validation error are both substantial but there is a little gap between them.

- Model too simple: unable to reduce the training error → **Underfitting**
- Model too complex: Training error « validation error → **Overfitting**
- Over- or underfitting depends on size of training data and model complexity

**Care more about validation error and find tradeoff**

# Underfitting or overfitting



## Example (Polynomial regression)

Given training data consisting of a single feature  $x$  and a corresponding real-valued label  $y$


- find the polynomial of degree  $M$

$$\begin{aligned}\hat{y} &= \sum_{j=0}^M x^j w_j \\ &= \sum_{j=0}^M \phi_j(x) w_j \\ &= \mathbf{w}^\top \mathbf{x} + b\end{aligned}$$

- Higher-order polynomial: More model parameters, lower training error

## Polynomial curve fitting


The degree  $M$  of the polynomial is crucial.



(C.M. Bishop, Pattern Recognition and Machine Learning)

# Polynomial curve fitting


The degree  $M$  of the polynomial is crucial.



(C.M. Bishop, Pattern Recognition and Machine Learning)

## Polynomial curve fitting

The degree  $M$  of the polynomial is crucial.




(C.M. Bishop, Pattern Recognition and Machine Learning)



# Polynomial curve fitting


The degree  $M$  of the polynomial is crucial.



(C.M. Bishop, Pattern Recognition and Machine Learning)

## Train vs. Test error


- Split sample in training and test set.
- Choose  $M$  based on test error.



(C.M. Bishop, Pattern Recognition and Machine Learning)



## $K$ -fold Cross validation:

- randomly assign samples  $(\mathbf{x}^{(i)}, y^{(i)})$  to  $K$  sets of equal size
- for each set  $s \in S$  (e.g.  $K = 4$  sets) and  $m \in [1, \dots, M]$ :
  - train model on  $K - 1$  remaining sets
  - predict on  $s$  and compute loss.
- compute average MSE for degree  $m$ .
- pick  $m$  with lowest loss.



(C.M. Bishop, Pattern Recognition and Machine Learning)

## Get more data



Others ways to avoid overfitting and fit complex model for limited number of observations?

## Regularization of weights

## Regularization

Regularize the regression weights.

Loss function:

$$\underbrace{\frac{1}{N} \sum_{i=1}^N \frac{1}{2} (\mathbf{w}^\top \mathbf{x}^{(i)} + b - y^{(i)})^2}_{l(\mathbf{w}, b)} + \underbrace{\frac{\lambda}{2} \sum_{j=1}^d w_j^2}_{l_2\text{-norm regularizer}}$$

where  $\sqrt{\sum_{j=1}^d w_j^2}$  is the  $l_2$ -norm of  $\mathbf{w}$ .



What effect does this have? How will the learned weights be different?

- Penalizes large weights.
- Reduces the complexity of the function that associates  $\mathbf{x}$  with  $\mathbf{y}$ , i.e. learn parsimonious model.
- Also known as **shrinkage** or **weight decay**.

## Regularization

Loss function:

$$\underbrace{\frac{1}{N} \sum_{i=1}^N \frac{1}{2} (\mathbf{w}^\top \mathbf{x}^{(i)} + b - y^{(i)})^2}_{l(\mathbf{w}, b)} + \underbrace{\frac{\lambda}{2} \sum_{j=1}^d w_j^2}_{l_2\text{-norm regularizer}}$$



The stochastic gradient descent updates for L2-regularised regression are as follows:

$$\mathbf{w} \leftarrow \left(1 - \frac{\eta\lambda}{|\mathcal{B}|}\right) \mathbf{w} - \frac{\eta}{|\mathcal{B}|} \sum_{i \in \mathcal{B}} \mathbf{x}^{(i)} \left( \mathbf{w}^\top \mathbf{x}^{(i)} + b - y^{(i)} \right),$$

# Weight decay




Figure: Training without regularization





Figure: Training with L2-regularization

## Regularization

$$\underbrace{\frac{1}{N} \sum_{i=1}^N \frac{1}{2} (\mathbf{w}^\top \mathbf{x}^{(i)} + b - y^{(i)})^2}_{l(\mathbf{w}, b)} + \underbrace{\frac{\lambda}{2} \sum_{j=1}^d w_j^2}_{l_2\text{-norm regularizer}}$$

Question: How to choose an optimal  $\lambda$ ?

Answer: Look at the test error!




(C.M. Bishop, Pattern Recognition and Machine Learning)



# Regularization

A more general regularization:

$$\underbrace{\frac{1}{N} \sum_{i=1}^N \frac{1}{2} (\mathbf{w}^\top \mathbf{x}^{(i)} + b - y^{(i)})^2}_{l(\mathbf{w}, b)} + \underbrace{\frac{\lambda}{2} \sum_{j=1}^d |w_j|^q}_{l_q\text{-norm regularizer}}$$



(C.M. Bishop, Pattern Recognition and Machine Learning)