

## 10.1. Attention Mechanism

Lecture based on “Dive into Deep Learning” <http://D2L.AI> (Zhang et al., 2020)

Prof. Dr. Christoph Lippert

Digital Health & Machine Learning

## Database queries

Databases are collections of keys ( $k$ ) and values ( $v$ ).

$\mathcal{D} = \{(\text{"Schumacher", "Anna"}, (\text{"Schneider", "Michael"}, (\text{"Firscher", "Julia"}, (\text{"Weber", "Sarah"}, (\text{"Schulz", "Lisa"}, (\text{"Schuler", "Tim"}))\}$ , with the last name being the key and the first name being the value.

- exact query (q) "Schulz" would return the value "Lisa".
- if ("Schulz", "Lisa") not in  $\mathcal{D}$ , there would be no valid answer.
- For approximate queries, we could retrieve ("Schuler", "Tim") instead.

- queries  $q$  operate on  $(k, v)$  pairs regardless of the database size.
- same query have different answers, according to the contents of the database.
- Simple queries executed on a large state space (the database) (e.g., exact match, approximate match, top- $k$ ).

The *attention mechanism* translates this concept to deep learning.

## Attention Pooling

$\mathcal{D} \stackrel{\text{def}}{=} \{(\mathbf{k}_1, \mathbf{v}_1), \dots, (\mathbf{k}_m, \mathbf{v}_m)\}$  a database of  $m$  keys and values.

Let  $\mathbf{q}$  be a query.

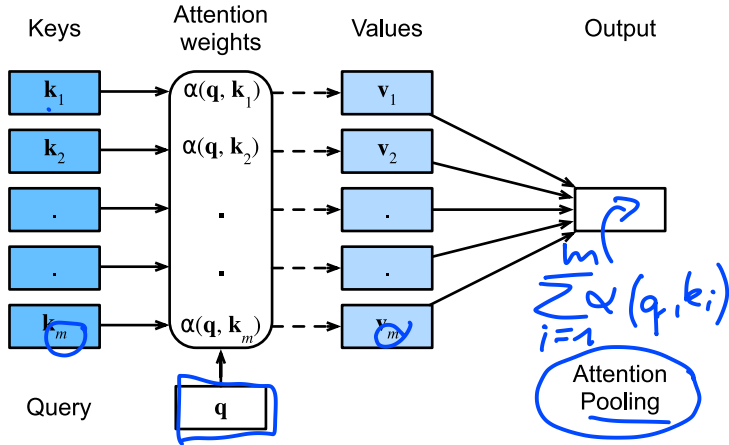
Then we can define the *attention* over  $\mathcal{D}$  as

$$\text{Attention}(\mathbf{q}, \mathcal{D}) \stackrel{\text{def}}{=} \sum_{i=1}^m \alpha(\mathbf{q}, \mathbf{k}_i) \mathbf{v}_i,$$


where  $\alpha(\mathbf{q}, \mathbf{k}_i) \in \mathbb{R}$  ( $i = 1, \dots, m$ ) are scalar attention weights.

- *attention* operation pays particular attention to the terms with large weight  $\alpha$ .
- Attention over  $\mathcal{D}$  generates a linear combination of values in the database.

# Attention Pooling



## Special Cases



- For nonnegative  $\alpha(\mathbf{q}, \mathbf{k}_i)$ , the output of the attention mechanism is contained in the convex cone spanned by the values  $\mathbf{v}_i$ .
- The weights  $\alpha(\mathbf{q}, \mathbf{k}_i)$  form a convex combination, i.e.,  $\sum_i \alpha(\mathbf{q}, \mathbf{k}_i) = 1$  and  $\alpha(\mathbf{q}, \mathbf{k}_i) \geq 0$  for all  $i$ .
- Exactly one of the weights  $\alpha(\mathbf{q}, \mathbf{k}_i)$  is 1, while all others are 0.
- If all weights are equal, i.e.,  $\alpha(\mathbf{q}, \mathbf{k}_i) = \frac{1}{m}$  for all  $i$ , the result is average pooling.

# Softmax Attention

Typically, attention weights are normalized to sum up to 1

$$\alpha(\mathbf{q}, \mathbf{k}_i) = \frac{\alpha(\mathbf{q}, \mathbf{k}_i)}{\sum_j \alpha(\mathbf{q}, \mathbf{k}_j)}$$

Any function  $a(\mathbf{q}, \mathbf{k})$  can be combined with softmax

$$\alpha(\mathbf{q}, \mathbf{k}_i) = \frac{\exp(a(\mathbf{q}, \mathbf{k}_i))}{\sum_j \exp(a(\mathbf{q}, \mathbf{k}_j))}$$


Benefits of softmax attention:

- available in all deep learning frameworks.
- differentiable
- gradient never vanishes

$$e^{-0.3}$$
$$\frac{1}{e^{0.3}}$$

- Alternatives possible
- non-differentiable attention can be trained using reinforcement learning.
- One of the benefits of the attention mechanism is that it can be quite intuitive, particularly when the weights are nonnegative and sum to 1.
- In this case we might *interpret* large weights as a way for the model to select components of relevance.
- While this is a good intuition, it is important to remember that it is just that, an *intuition*.

## Summary



- Attention allows us to aggregate data from many (key, value) pairs.
- Attention is a *differentiable* way by which a neural network can select elements from a set and to construct a weighted sum over representations.

Where do queries, keys, and values arise from in Machine Learning?

- in regression (Nadaraya-Watson estimator)
  - the query might correspond to the location where the regression should be carried out.
  - The values are the (regression) values themselves.
- Any of queries, keys and values are differentiable and can be learned from data in different Deep Learning architectures